

My execution times for the Sudoku solver ended up increasing with the number of cores, even with cancelling task groups to "prune" the search tree. This is most likely due to the overhead caused by creating the new threads and context switching which is made worse by the large amount of backtracking being done by the program.

Threads	Time (seconds)
1	1.28464
16	12.00283
32	19.10845
48	25.03550

Table 1: Execution Time vs. Thread Count