

Since regular vector addition doesn't reuse any values, I changed the kernel so that each element of the result is equal to the sum of all elements in the input arrays that are in the same block as the current thread, which made the kernel using shared memory noticeably outperform the kernel using global memory.

<b>Vector Size</b>	<b>Shared Memory Time (milliseconds)</b>	<b>Global Memory Time (milliseconds)</b>
1000	57	146
10000	94	273
100000	178	946
1000000	441	3588

Table 1: Execution Time vs. Vector Size